

# Optimizing Breadth-First Search on Modern Multicore CPUs

---

Salvatore D. Andaloro

Department of Information Engineering and Computer Science, University of Trento



UNIVERSITÀ  
DI TRENTO

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

**Introduction**

OpenMP

Pthreads

Results

Conclusions

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...



Social network

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

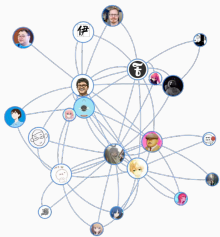
Pthreads

Results

Conclusions

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...



Social network



Road network

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

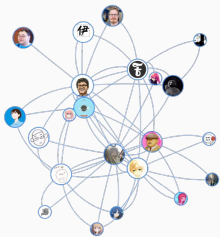
Pthreads

Results

Conclusions

# Breadth-First Search


- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...



Social network



Road network



Synthetic graph

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

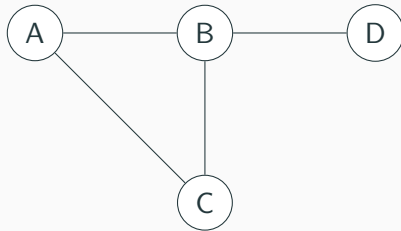
OpenMP

Pthreads

Results

Conclusions

# Breadth-First Search Example



**Source vertex: A**

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

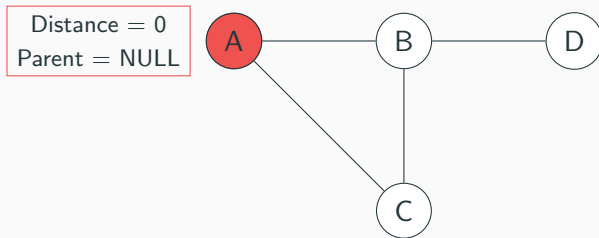
Pthreads

Results

Conclusions



# Breadth-First Search Example



**Frontier: A**

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

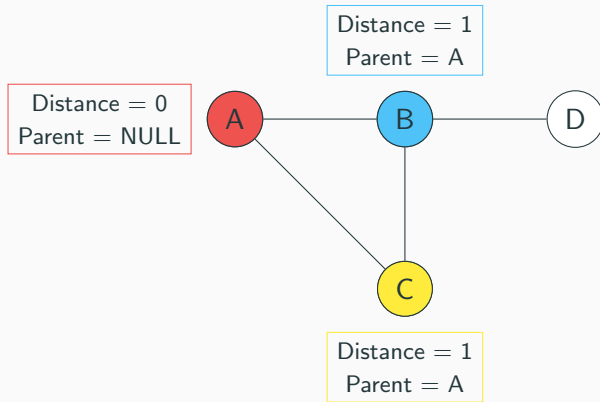
OpenMP

Pthreads

Results

Conclusions

# Breadth-First Search Example



**Frontier:** B, C

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

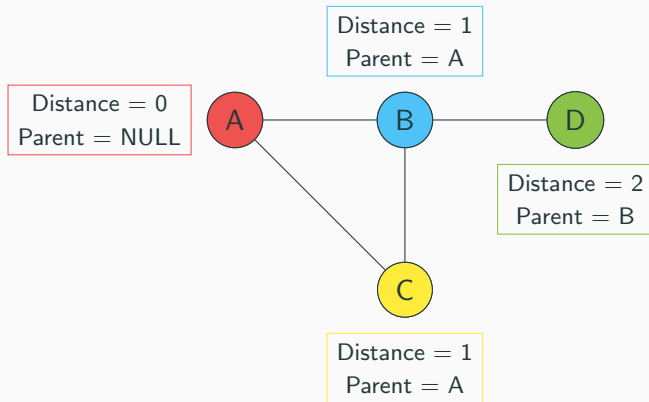
OpenMP

Pthreads

Results

Conclusions

# Breadth-First Search Example



**Frontier: D**

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
  - Cache effects must be considered

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
  - Cache effects must be considered
- CPUs exhibit growing amount of **parallelism...**



Evolution of core counts per socket for AMD and Intel processors

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
  - Cache effects must be considered
- CPUs exhibit growing amount of **parallelism**...
- ...and new architectures are coming to the market (ARM, RISC-V)



Evolution of core counts per socket for AMD and Intel processors

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Contents

- Two implementations with different **parallel programming paradigms**
  1. OpenMP implementation using the *MergedCSR* data structure
  2. Pthreads implementation using *MergedCSR* + custom synchronization routines

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions



# Contents

- Two implementations with different **parallel programming paradigms**
  1. OpenMP implementation using the *MergedCSR* data structure
  2. Pthreads implementation using *MergedCSR* + custom synchronization routines
- Evaluated against GAP Benchmark suite
- Speedups compared on three different architectures (AMD x86, RISC-V, ARM)



GAP suite logo



Compared architectures

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

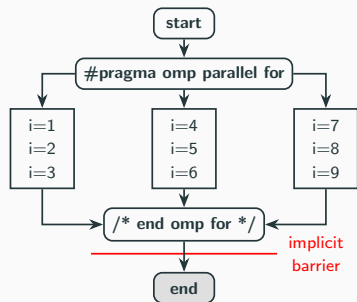
Results

Conclusions

# OpenMP implementation

- **OpenMP** is a widely used framework for **parallel programming** in C and C++
- Uses simple compiler directives called pragmas

```
#pragma omp parallel for  
for (int i = 1; i <= 9; i++) {  
    A[i] = i  
}
```



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

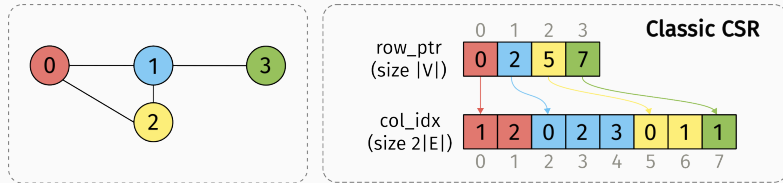
Pthreads

Results

Conclusions

# From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

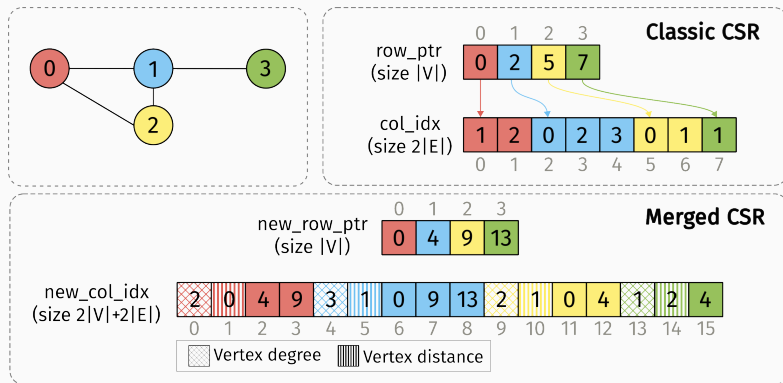
Pthreads

Results

Conclusions

# From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row\_ptr array during BFS traversal
  - row\_ptr array contains also algorithm-specific metadata (ex. distance)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

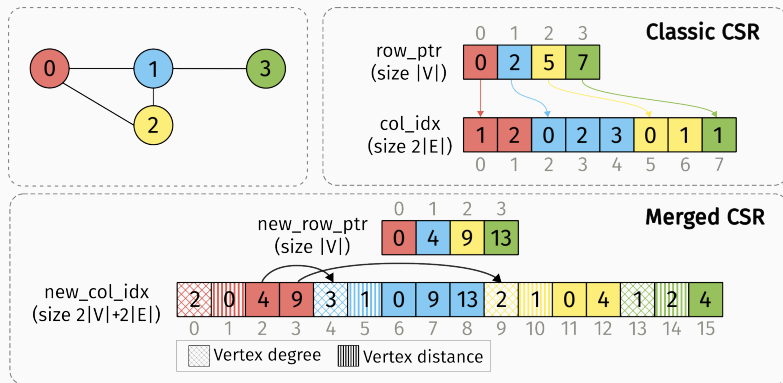
Pthreads

Results

Conclusions

# From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row\_ptr array during BFS traversal
  - row\_ptr array contains also algorithm-specific metadata (ex. distance)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs  
Salvatore D.  
Andaloro

Introduction

OpenMP

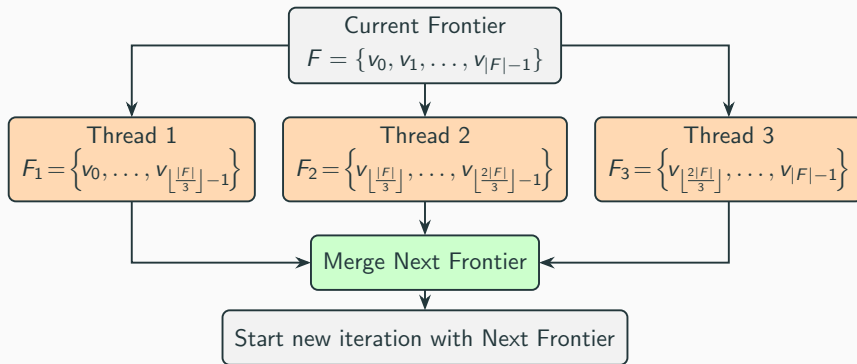
Pthreads

Results

Conclusions

# Parallelization strategies

- Different parallelization strategies, depending on the graph type
- Strategy used: Frontier partitioning + Merge step



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Implementation

```
#pragma omp declare reduction(vec_add : \
    omp_out.insert(omp_out.end(), omp_in.begin(), omp_in.end()))

#pragma omp parallel for reduction(vec_add : next_frontier)
↪ if(this_frontier.size() > 50)
for (const auto &v : this_frontier) {
    for (vertex i = v + 2; i < end; i++) { // Iterate over neighbors
        vertex neighbor = new_col_idx[i];
        // If neighbor is not visited, add to frontier
        if (DISTANCE(neighbor) == max()) {
            next_frontier.push_back(neighbor);
            DISTANCE(neighbor) = distance; // Set the distance
        }
    }
}
```

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Inefficiencies of the OpenMP implementation

- Merging step is not parallel

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

**OpenMP**

Pthreads

Results

Conclusions



# Inefficiencies of the OpenMP implementation

- Merging step is not parallel
- Poor cache locality, as vertices are collected and repartitioned among the cores

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Inefficiencies of the OpenMP implementation

- Merging step is not parallel
- Poor cache locality, as vertices are collected and repartitioned among the cores
- For **large-diameter graphs**, OpenMP enters the `parallel` region more than 10k times for a single BFS runs

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C



Pthreads (unofficial)  
logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

**Pthreads**

Results

Conclusions

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier



Pthreads (unofficial)  
logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

**Pthreads**

Results

Conclusions

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier
  2. Work-stealing mechanism for load balancing



Pthreads (unofficial)  
logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

**Pthreads**

Results

Conclusions

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier
  2. Work-stealing mechanism for load balancing
  3. Thread pool to manage thread creation and destruction



Pthreads (unofficial)  
logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier
  2. Work-stealing mechanism for load balancing
  3. Thread pool to manage thread creation and destruction
  4. Custom barrier for thread synchronization



Pthreads (unofficial)  
logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

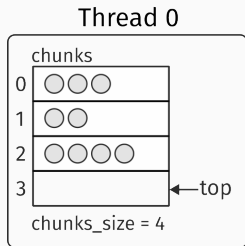
OpenMP

**Pthreads**

Results

Conclusions

# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions



# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

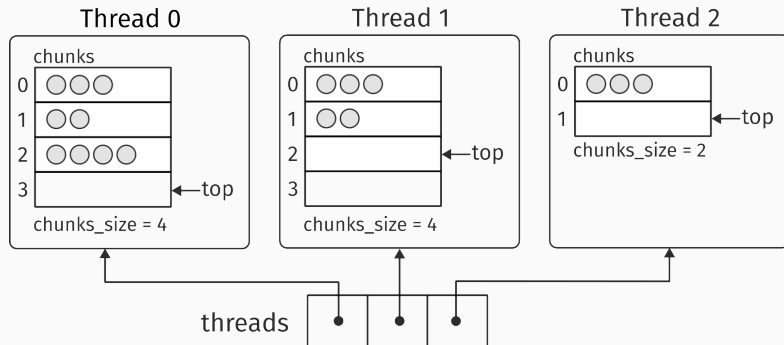
OpenMP

Pthreads

Results

Conclusions

# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

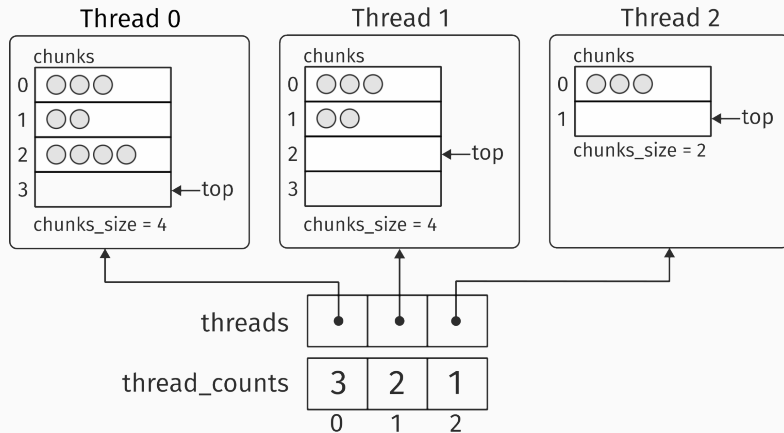
OpenMP

Pthreads

Results

Conclusions

# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

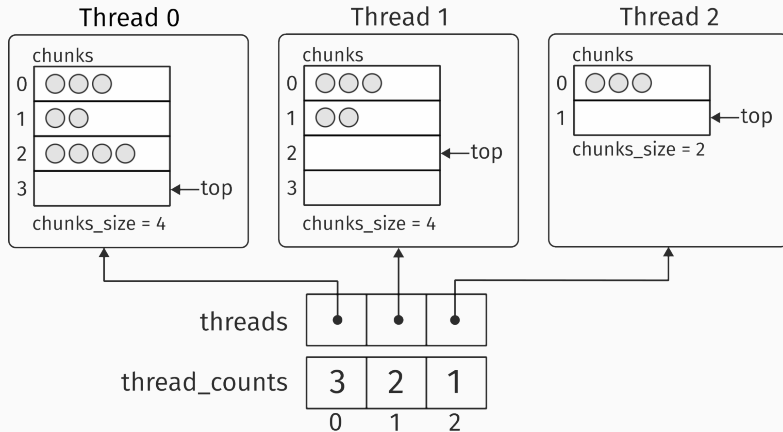
Pthreads

Results

Conclusions

# Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

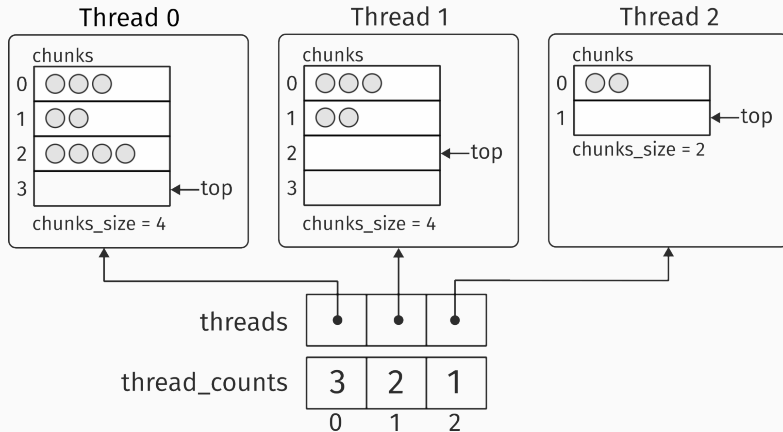
Pthreads

Results

Conclusions

# Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

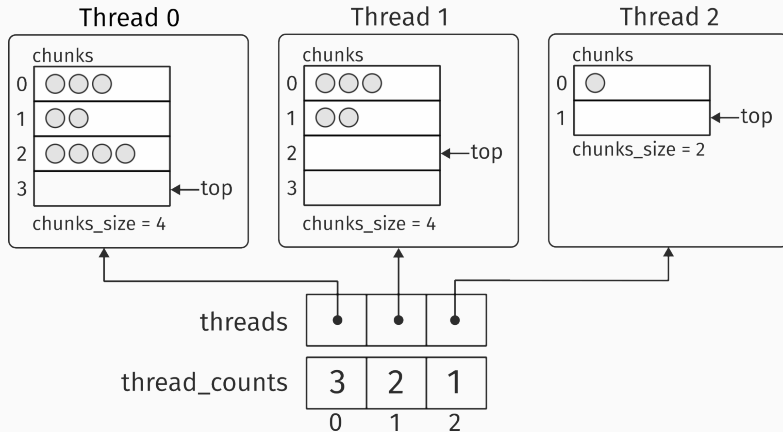
Pthreads

Results

Conclusions

# Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

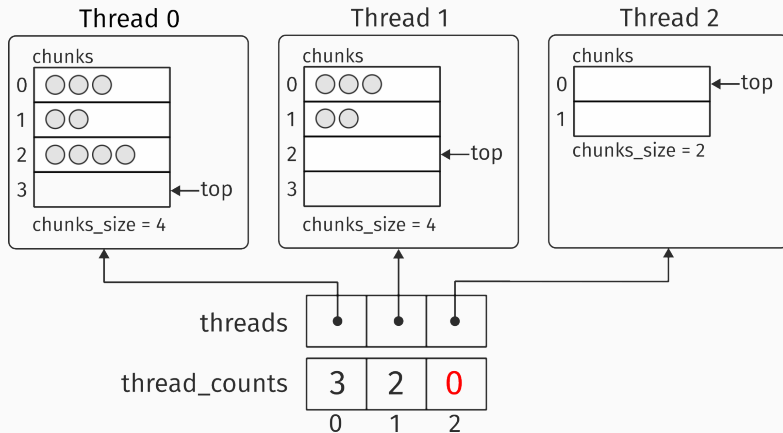
Pthreads

Results

Conclusions

# Work-stealing mechanism

Thread 2 is out of work, will attempt a steal soon...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

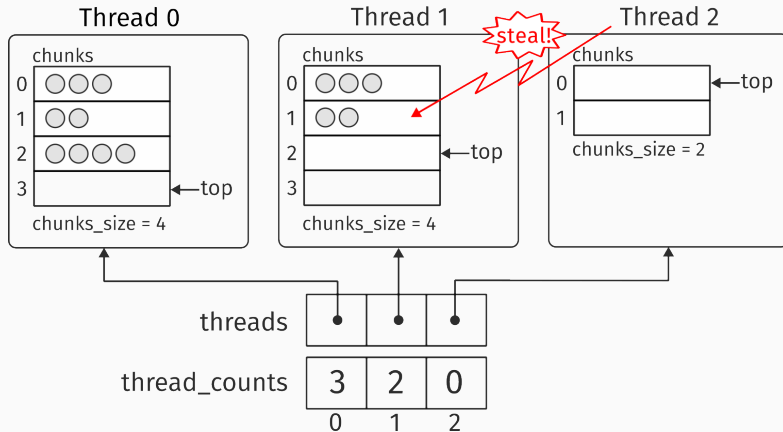
Pthreads

Results

Conclusions

# Work-stealing mechanism

Thread 2 steals a chunk of work from Thread 1...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

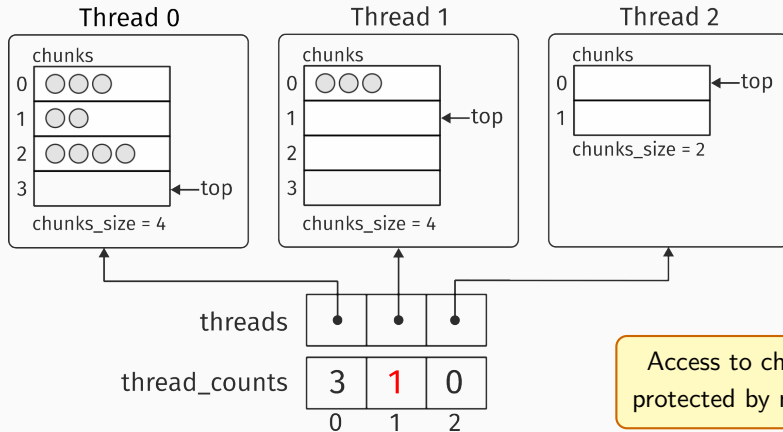
Results

Conclusions



# Work-stealing mechanism

Thread 2 processes the stolen vertices and updates the global count.



Access to chunks is protected by mutexes.

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions


# Thread pool

- When the program is run, a group of threads is spawned
- At the beginning of each BFS run, the threads are awakened and the starting vertex is assigned to the 0<sup>th</sup> thread



# Sense-Reversal Centralized Barrier

- Barrier: point that threads must reach before any can proceed
- Procedure:
  1. Central counter tracks arriving threads
  2. Last thread resets counter + increment distance
  3. Others threads spin wait until distance changes
  4. All threads are released together



# Experimental setup

- Experiments run on 3 platforms:
  - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
  - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
  - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Experimental setup

- Experiments run on 3 platforms:
  - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
  - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
  - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)
- Datasets: 3 road networks (USA, Europe, Asia), 3 FEM meshes (Earth's crust, steel hook, porous material), 1 random geometric graph (RGG)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Experimental setup

- Experiments run on 3 platforms:
  - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
  - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
  - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)
- Datasets: 3 road networks (USA, Europe, Asia), 3 FEM meshes (Earth's crust, steel hook, porous material), 1 random geometric graph (RGG)
- Tools: GCC compiler, Likwid, SBatchMan



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Experimental setup

- Experiments run on 3 platforms:
  - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
  - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
  - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)
- Datasets: 3 road networks (USA, Europe, Asia), 3 FEM meshes (Earth's crust, steel hook, porous material), 1 random geometric graph (RGG)
- Tools: GCC compiler, Likwid, SBatchMan
- Compared against the GAP benchmark suite



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

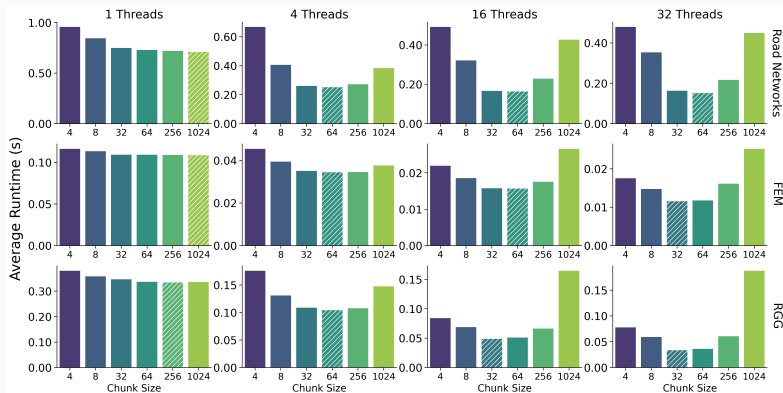
Pthreads

Results

Conclusions

# Chunk size impact on performance

- Chunk size determines the number of vertices in a chunk
- Chunk sizes of 32 and 64 are optimal for most datasets in multithreaded environments



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

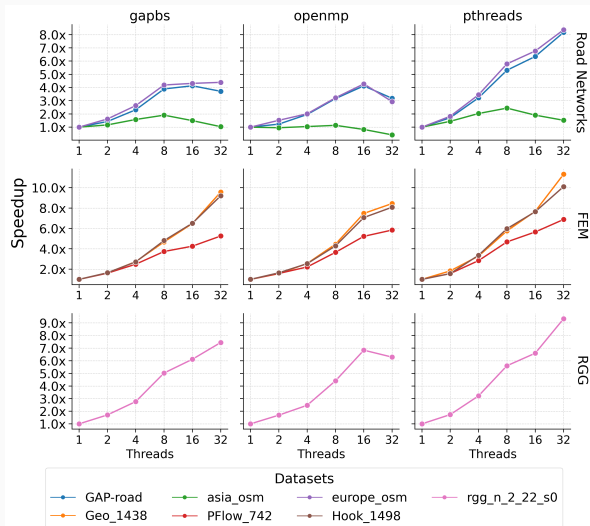
Pthreads

Results

Conclusions



# Scalability



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

# Speedup - OpenMP



Speedup of the OpenMP implementation compared to the GAPBS implementation

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

# Speedup - Pthreads



Speedup of the pthreads implementation compared to the OpenMP implementation

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

# Comparison on different architectures



Execution time and speedup on different architectures for the Europe road network dataset

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs  
  
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Conclusions

- Presented a multithreaded implementation of the BFS algorithm using OpenMP and Pthreads
- Compared it on different architectures (x86, RISC-V, ARM) and different datasets
- Achieved  $\approx 1.5x$  geomean speedup for OpenMP and  $\approx 2x$  speedup for Pthreads compared to the GAP benchmark suite
- Future work: explore other graph algorithms, optimize for more graph types, use different barrier or synchronization primitives

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

# Thank You!

Questions?