

# Optimizing Breadth-First Search on Modern Multicore CPUs

---

Salvatore D. Andaloro

Department of Information Engineering and Computer Science, University of Trento



UNIVERSITÀ  
DI TRENTO

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

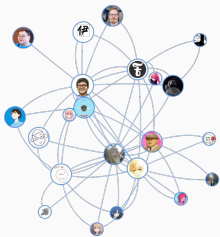
Introduction

OpenMP

Pthreads

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...



Social network

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

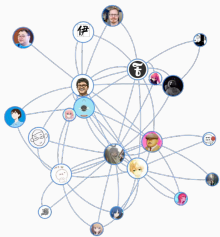
Introduction

OpenMP

Pthreads

# Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...



Social network



Road network

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

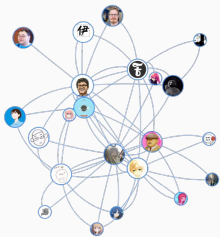
Introduction

OpenMP

Pthreads

# Breadth-First Search

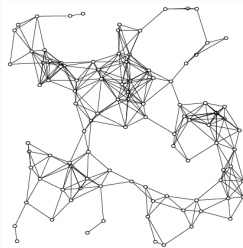
- Breadth-First Search is a fundamental algorithm in graph analysis
- Vertices are labeled based on the **distance** from a given *source* vertex
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...



Social network



Road network



Synthetic graph

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

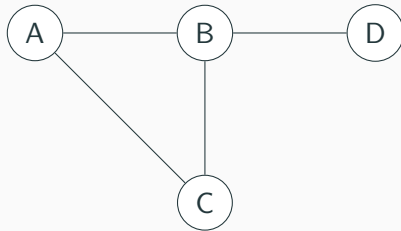
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Breadth-First Search Example



**Source vertex: A**

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

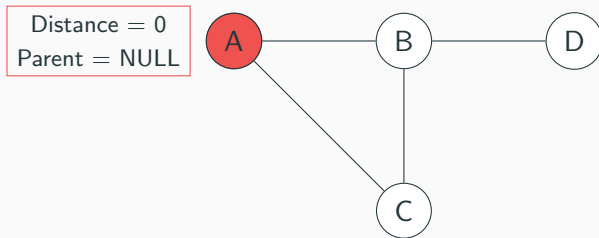
Introduction

OpenMP

Pthreads



# Breadth-First Search Example



**Frontier: A**

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

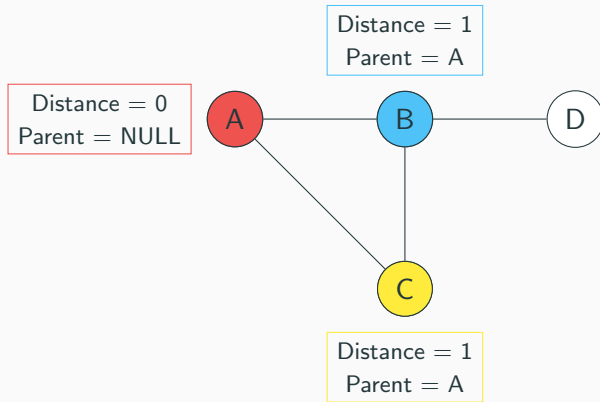
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Breadth-First Search Example



**Frontier:** B, C

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

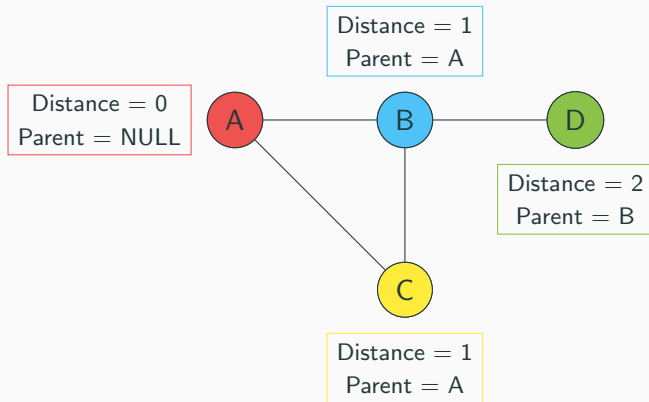
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Breadth-First Search Example



**Frontier:** D

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
  - Cache effects must be considered

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

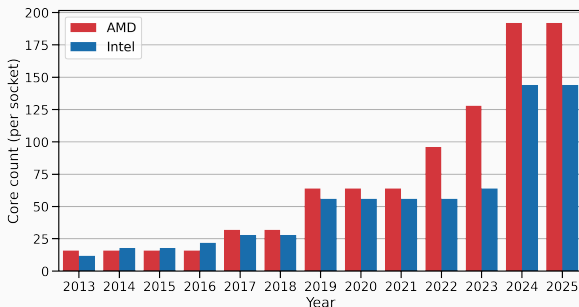
Introduction

OpenMP

Pthreads

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
  - Cache effects must be considered
- CPUs exhibit growing amount of **parallelism...**



Evolution of core counts per socket for AMD and Intel processors

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

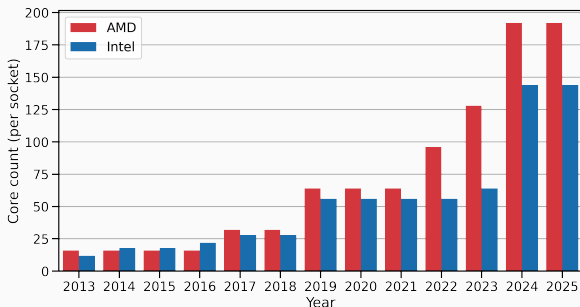
Introduction

OpenMP

Pthreads

# Modern Computer Architectures

- BFS has  $\mathcal{O}(V + E)$  time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
  - Cache effects must be considered
- CPUs exhibit growing amount of **parallelism**...
- ...and new architectures are coming to the market (ARM, RISC-V)



Evolution of core counts per socket for AMD and Intel processors

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Contents

- Two implementations with different **parallel programming paradigms**
  1. OpenMP implementation using the *MergedCSR* data structure
  2. Pthreads implementation using *MergedCSR* + custom synchronization routines

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

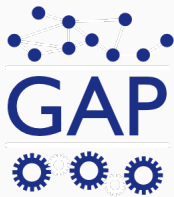
OpenMP

Pthreads



# Contents

- Two implementations with different **parallel programming paradigms**
  1. OpenMP implementation using the *MergedCSR* data structure
  2. Pthreads implementation using *MergedCSR* + custom synchronization routines
- Evaluated against GAP Benchmark suite
- Speedups compared on three different architectures (AMD x86, RISC-V, ARM)



GAP suite logo



Compared architectures

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

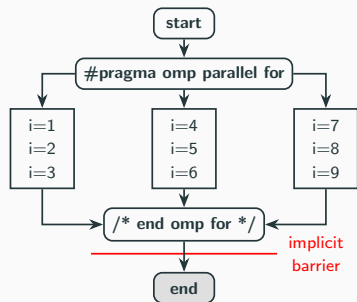
OpenMP

Pthreads

# OpenMP implementation

- **OpenMP** is a widely used framework for **parallel programming** in C and C++
- Uses simple compiler directives called pragmas

```
#pragma omp parallel for  
for (int i = 1; i <= 9; i++) {  
    A[i] = i  
}
```



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

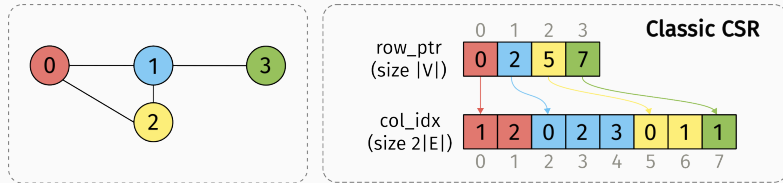
Introduction

OpenMP

Pthreads

# From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

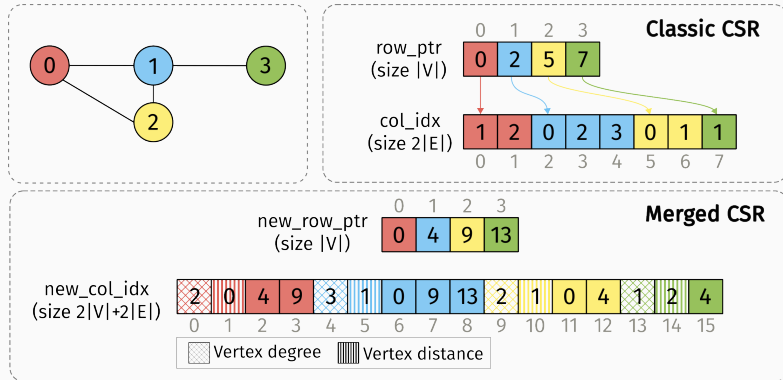
Introduction

OpenMP

Pthreads

# From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row\_ptr array during BFS traversal
  - row\_ptr array contains also algorithm-specific metadata (ex. distance)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

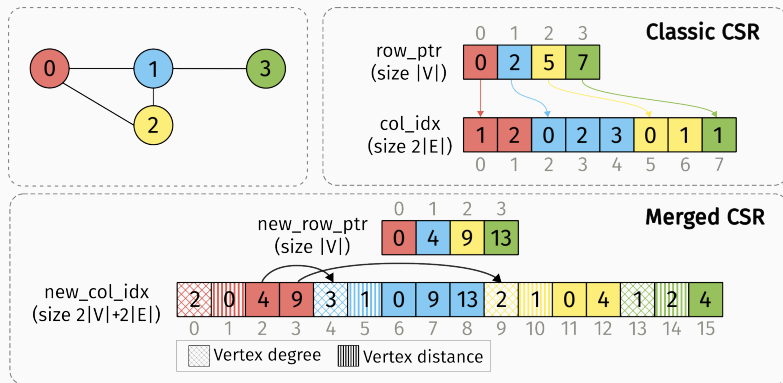
Introduction

OpenMP

Pthreads

# From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row\_ptr array during BFS traversal
  - row\_ptr array contains also algorithm-specific metadata (ex. distance)



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

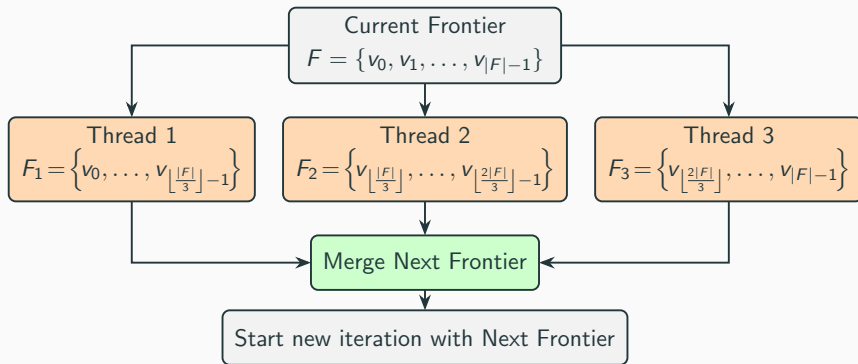
Introduction

OpenMP

Pthreads

# Parallelization strategies

- Different parallelization strategies, depending on the graph type
- Strategy used: Frontier partitioning + Merge step



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Implementation

```
#pragma omp declare reduction(vec_add : \
    omp_out.insert(omp_out.end(), omp_in.begin(), omp_in.end()))

#pragma omp parallel for reduction(vec_add : next_frontier)
↪ if(this_frontier.size() > 50)
for (const auto &v : this_frontier) {
    for (vertex i = v + 2; i < end; i++) { // Iterate over neighbors
        vertex neighbor = new_col_idx[i];
        // If neighbor is not visited, add to frontier
        if (DISTANCE(neighbor) == max()) {
            next_frontier.push_back(neighbor);
            DISTANCE(neighbor) = distance; // Set the distance
        }
    }
}
```

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Inefficiencies of the OpenMP implementation

- Merging step is not parallel

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads



# Inefficiencies of the OpenMP implementation

- Merging step is not parallel
- Poor cache locality, as vertices are collected and repartitioned among the cores

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Inefficiencies of the OpenMP implementation

- Merging step is not parallel
- Poor cache locality, as vertices are collected and repartitioned among the cores
- For **large-diameter graphs**, OpenMP enters the `parallel` region more than 10k times for a single BFS runs

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C



Pthreads (unofficial) logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier



Pthreads (unofficial) logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier
  2. Work-stealing mechanism for load balancing



Pthreads (unofficial) logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier
  2. Work-stealing mechanism for load balancing
  3. Thread pool to manage thread creation and destruction



Pthreads (unofficial) logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
  1. Custom data structure to handle the vertices in the frontier
  2. Work-stealing mechanism for load balancing
  3. Thread pool to manage thread creation and destruction
  4. Custom barrier for thread synchronization



Pthreads (unofficial) logo

Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

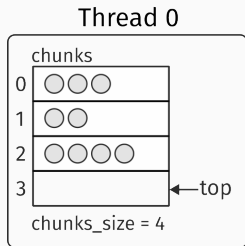
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

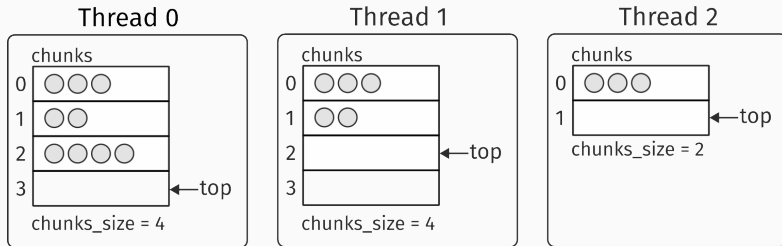
Introduction

OpenMP

Pthreads



# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

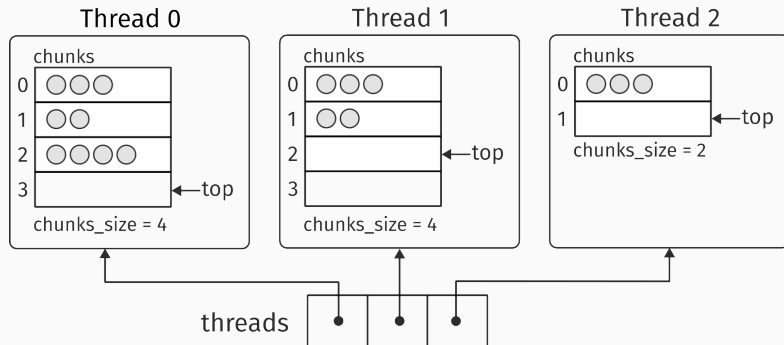
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

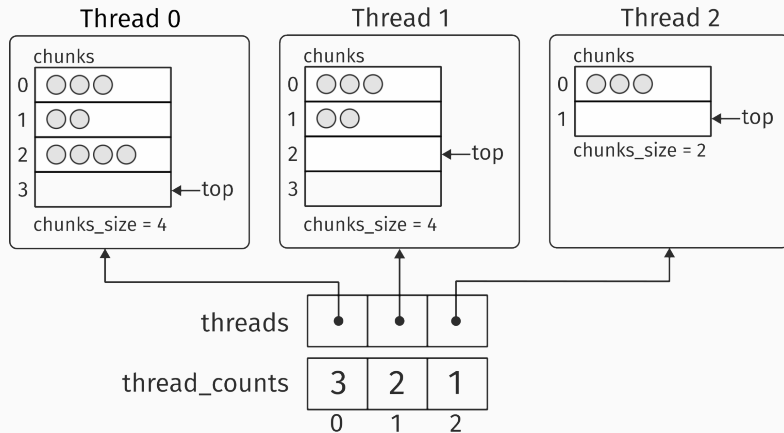
Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads

# Frontier implementation



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

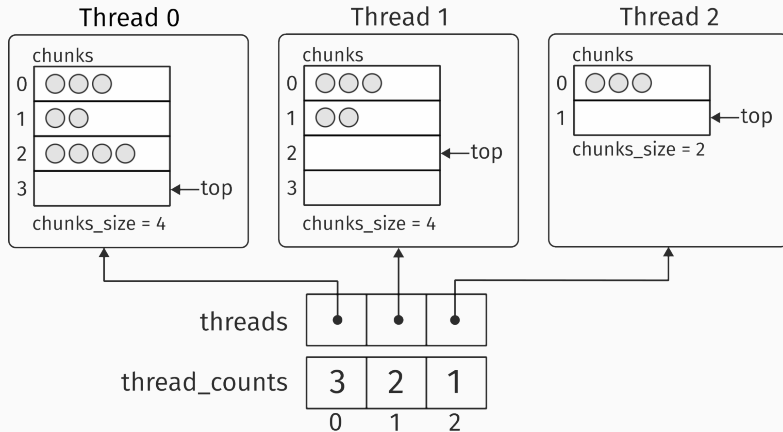
Introduction

OpenMP

Pthreads

# Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

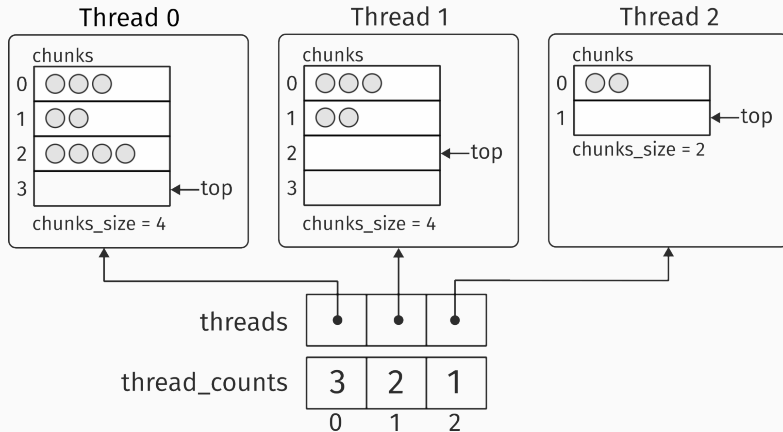
Introduction

OpenMP

Pthreads

# Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

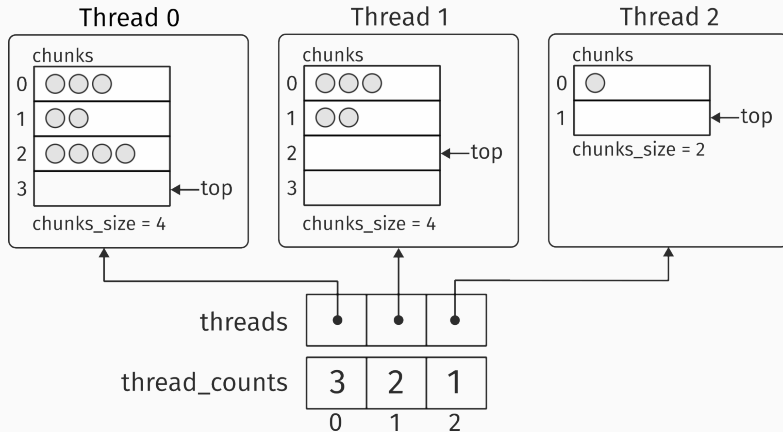
Introduction

OpenMP

Pthreads

# Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

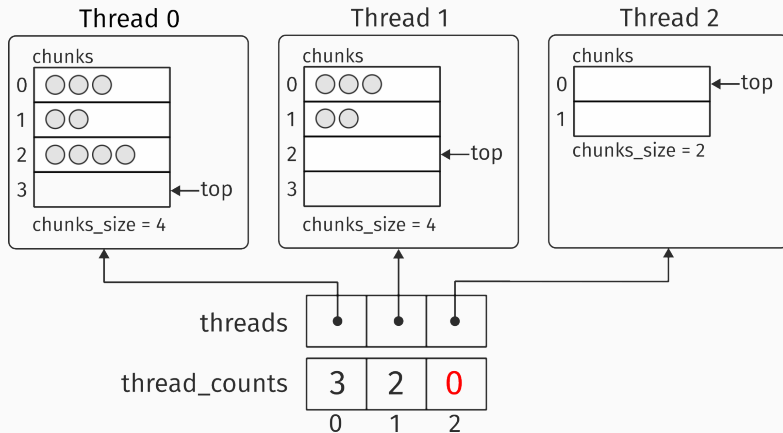
Introduction

OpenMP

Pthreads

# Work-stealing mechanism

Thread 2 is out of work, will attempt a steal soon...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

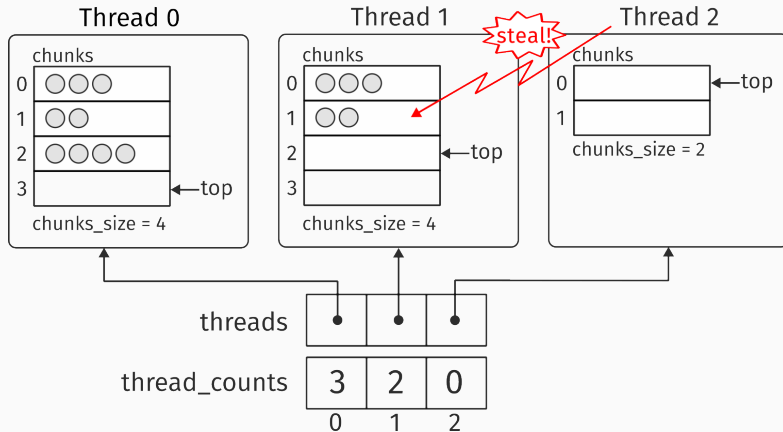
Introduction

OpenMP

Pthreads

# Work-stealing mechanism

Thread 2 steals a chunk of work from Thread 1...



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

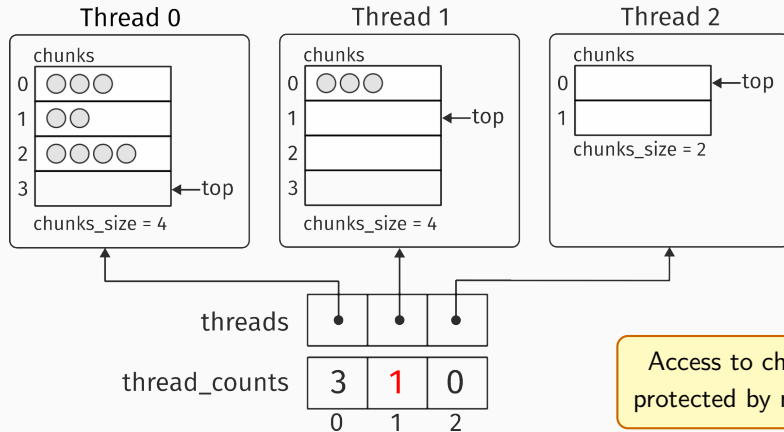
OpenMP

Pthreads



# Work-stealing mechanism

Thread 2 processes the stolen vertices and updates the global count.



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

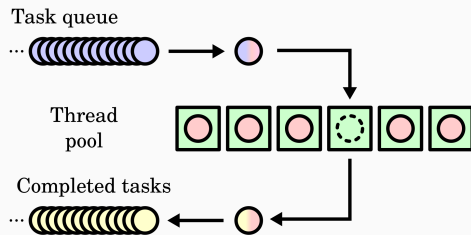
Introduction

OpenMP

Pthreads

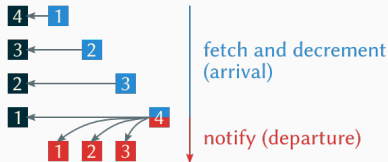
# Thread pool

- When the program is run, a group of threads is spawned
- At the beginning of each BFS run, the threads are awoken and the starting vertex is assigned to the 0<sup>th</sup> thread



# Sense-Reversal Centralized Barrier

- Central counter tracks arriving threads
- Last thread resets counter + toggles global sense
- Others wait until global = local sense
- All released together, barrier reusable



Optimizing  
Breadth-First  
Search on Modern  
Multicore CPUs

Salvatore D.  
Andaloro

Introduction

OpenMP

Pthreads