

Optimizing Breadth-First Search on Modern Multicore CPUs


Salvatore D. Andaloro

Department of Information Engineering and Computer Science, University of Trento



Breadth-First Search

- Breadth-First Search is a fundamental algorithm in graph analysis
- Used in many algorithms: Dijkstra, Maximum Flow, MSP...
- Vertices are labeled based on the **distance** from a given *source* vertex



Social network



Road network

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Breadth-First Search Example



Source vertex: A

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Breadth-First Search Example



Frontier: A

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

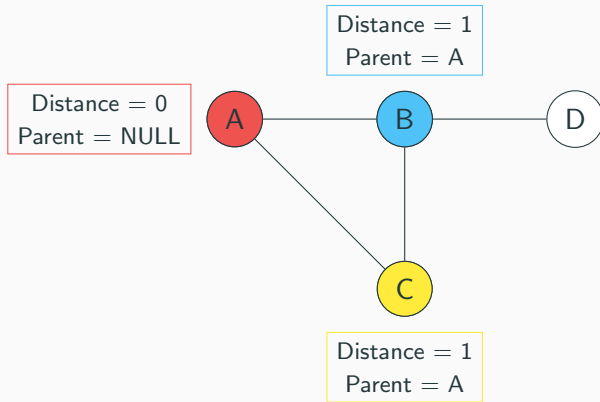
OpenMP

Pthreads

Results

Conclusions

Breadth-First Search Example



Frontier: B, C

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

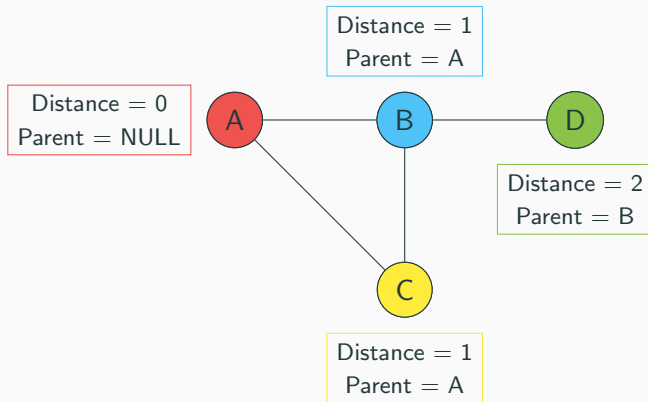
OpenMP

Pthreads

Results

Conclusions

Breadth-First Search Example



Frontier: D

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Modern Computer Architectures

- BFS has $\mathcal{O}(V + E)$ time and space complexity (under RAM model)

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Modern Computer Architectures

- BFS has $\mathcal{O}(V + E)$ time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
 - Algorithm exhibits poor cache locality

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

Modern Computer Architectures

- BFS has $\mathcal{O}(V + E)$ time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
 - Algorithm exhibits poor cache locality
- CPUs exhibit growing amount of **parallelism...**



Evolution of core counts per socket for AMD and Intel processors

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

Modern Computer Architectures

- BFS has $\mathcal{O}(V + E)$ time and space complexity (under RAM model)
- In practice, it is a **memory-bound algorithm**
 - Algorithm exhibits poor cache locality
- CPUs exhibit growing amount of **parallelism**...
- ...and new architectures are coming to the market (ARM, RISC-V)



Evolution of core counts per socket for AMD and Intel processors

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Contents

- New *MergedCSR* data structure
- Two optimized parallel implementations (OpenMP and pthreads)

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Contents

- New *MergedCSR* data structure
- Two optimized parallel implementations (OpenMP and pthreads)
- Evaluated against GAP Benchmark suite
- Speedups compared on three different architectures (AMD x86, RISC-V, ARM)



GAP suite logo



Compared architectures

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

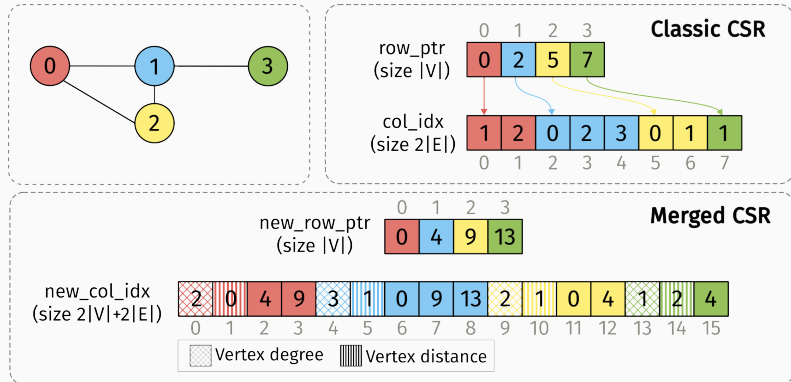
Pthreads

Results

Conclusions

From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row_ptr array during BFS traversal
 - row_ptr array contains also algorithm-specific metadata (ex. distance)



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

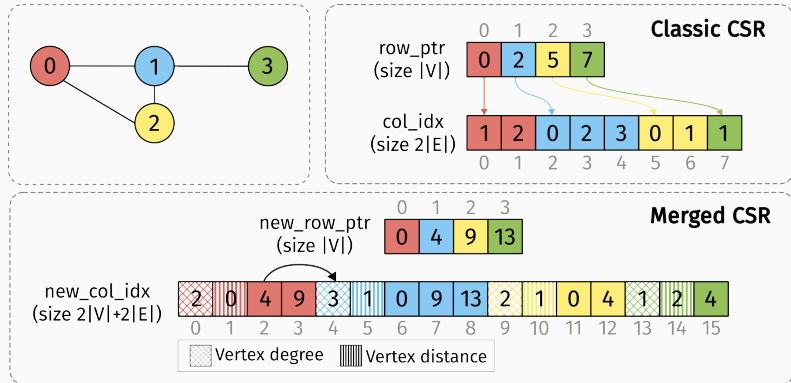
Pthreads

Results

Conclusions

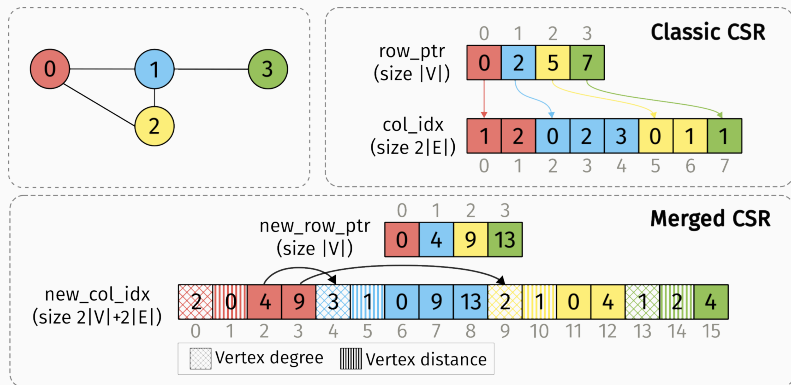
From CSR to MergedCSR

- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row_ptr array during BFS traversal
 - row_ptr array contains also algorithm-specific metadata (ex. distance)



From CSR to MergedCSR

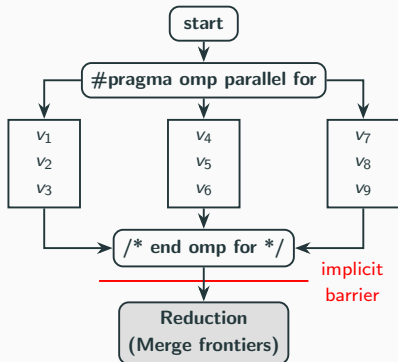
- Graphs are usually stored in the Compressed Sparse Row format (CSR)
- MergedCSR core idea: access only row_ptr array during BFS traversal
 - row_ptr array contains also algorithm-specific metadata (ex. distance)



OpenMP implementation

```
#pragma omp parallel for  
↪ reduction(vec_add : next_frontier)  
for (const auto &v : this_frontier) {  
    // Process vertex v  
}
```

- **OpenMP** is a widely used framework for **parallel programming** in C and C++
- Uses simple compiler directives called pragmas



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Inefficiencies of the OpenMP implementation

- Merging step is not parallel

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Inefficiencies of the OpenMP implementation

- Merging step is not parallel
- In graphs with large diameter, OpenMP enters the `parallel` for region more than 10k times for a single BFS runs

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C



Pthreads (unofficial)
logo

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
 1. Custom data structure to handle the vertices in the frontier



Pthreads (unofficial)
logo

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
 1. Custom data structure to handle the vertices in the frontier
 2. Work-stealing mechanism for load balancing



Pthreads (unofficial)
logo

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
 1. Custom data structure to handle the vertices in the frontier
 2. Work-stealing mechanism for load balancing
 3. Thread pool to manage thread creation and destruction



Pthreads (unofficial)
logo

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Pthreads implementation summary

- Pthreads: low-level threading library to create and manage threads in C
- Implementation components:
 1. Custom data structure to handle the vertices in the frontier
 2. Work-stealing mechanism for load balancing
 3. Thread pool to manage thread creation and destruction
 4. Custom barrier for thread synchronization



Pthreads (unofficial)
logo

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Frontier implementation



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Frontier implementation



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Frontier implementation



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Frontier implementation



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

Work-stealing mechanism

Thread 2 processes its vertices...



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions

Work-stealing mechanism

Thread 2 is out of work, will attempt a steal soon...



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

[Introduction](#)

[OpenMP](#)


[Pthreads](#)

[Results](#)

[Conclusions](#)

Work-stealing mechanism

Thread 2 steals a chunk of work from Thread 1...



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

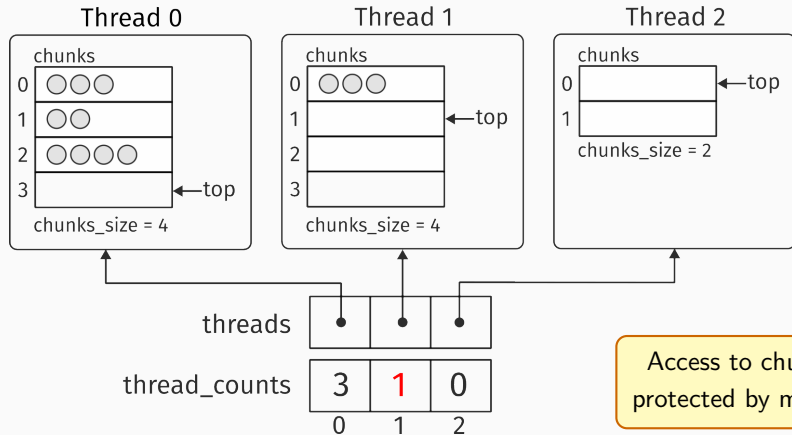
Pthreads

Results

Conclusions

Work-stealing mechanism

Thread 2 processes the stolen vertices and updates the global count.



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP


Pthreads

Results

Conclusions


Thread pool

- When the program is run, a group of threads is spawned
- At the beginning of each BFS run, the threads are awoken
 1. Process own chunks
 2. Steal work from other threads



Sense-Reversal Centralized Barrier

- Barrier: point that threads must reach before any can proceed
- Procedure:
 1. Central counter tracks arriving threads
 2. Last thread resets counter + increment distance
 3. Others threads spin wait until distance changes
 4. All threads are released together



Experimental setup

- Experiments run on 3 platforms:
 - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
 - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
 - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Experimental setup

- Experiments run on 3 platforms:
 - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
 - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
 - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)
- Datasets: 3 road networks (USA, Europe, Asia), 3 FEM meshes (Earth's crust, steel hook, porous material), 1 random geometric graph (RGG)



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Experimental setup

- Experiments run on 3 platforms:
 - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
 - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
 - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)
- Datasets: 3 road networks (USA, Europe, Asia), 3 FEM meshes (Earth's crust, steel hook, porous material), 1 random geometric graph (RGG)
- Tools: GCC compiler, Likwid, SBatchMan



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Experimental setup

- Experiments run on 3 platforms:
 - AMD EPYC 7543 CPU @ 2.8 GHz (32 cores)
 - Sophon SG2042 RISC-V CPU @ 2.0 GHz (64 cores)
 - NVIDIA Grace CPU Superchip @ up to 3.0 GHz (144 cores)
- Datasets: 3 road networks (USA, Europe, Asia), 3 FEM meshes (Earth's crust, steel hook, porous material), 1 random geometric graph (RGG)
- Tools: GCC compiler, Likwid, SBatchMan
- Compared against the GAP benchmark suite



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

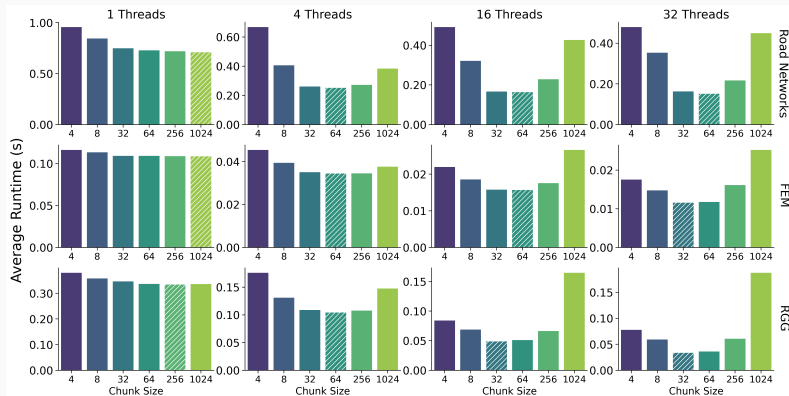
Pthreads

Results

Conclusions

Chunk size impact on performance

- Chunk size determines the number of vertices in a chunk
- Chunk sizes of 32 and 64 are optimal for most datasets in multithreaded environments



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Scalability



Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Speedup - OpenMP



Speedup of the OpenMP implementation compared to the GAPBS implementation

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Speedup - Pthreads



Speedup of the pthreads implementation compared to the OpenMP implementation

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction


OpenMP

Pthreads

Results

Conclusions

Comparison on different architectures



Execution time and speedup on different architectures for the Europe road network dataset

Optimizing
Breadth-First
Search on Modern
Multicore CPUs
Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Conclusions

- Presented a multithreaded implementation of the BFS algorithm using OpenMP and Pthreads
- Compared it on different architectures (x86, RISC-V, ARM) and different datasets
- Achieved $\approx 1.5x$ geomean speedup for OpenMP and $\approx 2x$ speedup for Pthreads compared to the GAP benchmark suite
- Future work: explore other graph algorithms, optimize for more graph types, use different barrier or synchronization primitives

Optimizing
Breadth-First
Search on Modern
Multicore CPUs

Salvatore D.
Andaloro

Introduction

OpenMP

Pthreads

Results

Conclusions

Thank You!

Questions?